

CUDA Tutorial

COMP5112 Assignment3

Outline

- CUDA Environment
- CUDA programming basics
- Assignment 3

CUDA Environment on CS Lab2

- CUDA version: 8.0
 - path: /usr/local/cuda-8.0/
- Check your CUDA environment first in the terminal:
 - Use `nvcc --version`
- If you cannot find `nvcc` command (or it is not CUDA 8.0), please add the CUDA toolkit installation path to the end of your `~/.cshrc_user` file

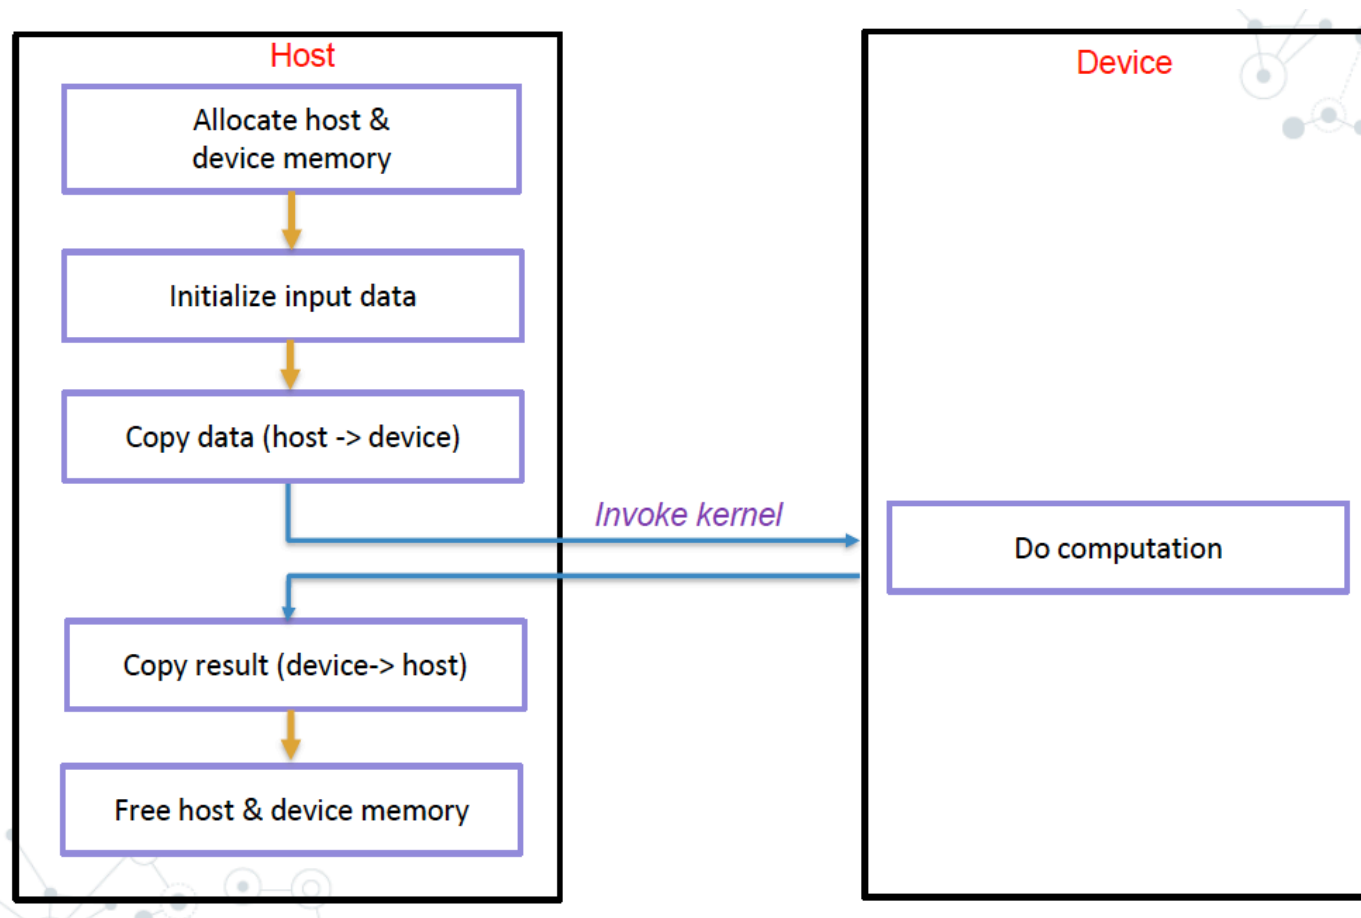
```
csl2wk01:ylwanghz:156> nvcc --version
nvcc: NVIDIA (R) Cuda compiler driver
Copyright (c) 2005-2016 NVIDIA Corporation
Built on Tue_Jan_10_13:22:03_CST_2017
Cuda compilation tools, release 8.0, V8.0.61
csl2wk01:ylwanghz:157> █
```

```
#####
#      File : .cshrc_user
#      Generic Version 1.1. CS. HKUST
#####

setenv MANPATH /usr/local/share/man:${MANPATH}
setenv PATH "${PATH}:/usr/local/software/openmpi/bin"
setenv PATH "${PATH}:/usr/local/cuda-8.0/bin/"
#export CUDA_DEBUGGER_SOFTWARE_PREEMPTION=1
set cuda software_preemption on

~
```

Typical CUDA programming model



Memory Allocation

- Host Memory

- malloc
 - `void* malloc(size_t size);`
- Parameters:
 - `size`: size of the memory block, in bytes. `size_t` is an unsigned integral type.
- returns:
 - On success, a pointer to the memory block allocated by the function.

- Device Memory

- cudaMalloc
 - `cudaMalloc(void **ptr, size_t size);`
- Parameters:
 - `**ptr`: Pointer to allocated device memory.
 - `size`: Requested allocation size in bytes.
- returns:
 - `cudaSuccess`, `cudaErrorMemoryAllocation`

Example

```
int *h_A, *d_A;  
size_t size = 1024* sizeof(int);
```

```
//on host memory  
h_A = (int*) malloc(size);
```

```
//on device memory  
cudaMalloc(&d_A, size);
```

Memory deallocation

- Host Memory

- free
 - `void* free(void *ptr);`
- Parameters:
 - `*ptr`: This is the pointer to a memory block previously allocated with `malloc`, `calloc` or `realloc` to be deallocated. If a null pointer is passed as argument, no action occurs.
- returns:
 - This function does not return any value.

- Device Memory

- `cudaFree`
 - `cudaMalloc(void **ptr);`
- Parameters:
 - `**ptr`: Device pointer to memory to free.
- returns:
 - `cudaSuccess`,
`cudaErrorInvalidDevicePointer`,
`cudaErrorInitializationError`.

Example

```
int *h_A, *d_A;  
size_t size = 1024* sizeof(int);
```

```
//allocate memory  
h_A = (int*) malloc(size);  
cudaMalloc(&d_A, size);
```

```
//free memory on host  
free(h_A);
```

```
//free memory on device  
cudaFree(d_A);
```

Data transfer between host and device

```
cudaMemcpy(void*      dst,  
           const void* src,  
           size_t     count,  
           enum cudaMemcpyKind kind)
```

kind:

```
cudaMemcpyHostToHost,  
cudaMemcpyHostToDevice,  
cudaMemcpyDeviceToHost,  
cudaMemcpyDeviceToDevice
```

//host -> device

cudaMemcpy(d_A, h_A, size, cudaMemcpyHostToDevice)

//device -> host

cudaMemcpy(h_A, d_A, size, cudaMemcpyDeviceToHost)

CUDA Kernel Declaration and Invocation

- A kernel function declaration has the prefix `__global__`, return type `void`.

```
__global__ void kernelName(param1, ..);
```

- A kernel function invocation includes launch parameters: `#block`, `#thread`..

```
kernelName<<<#block, #thread, size, stream>>>(param1, ..);
```

`#block`: number of blocks per grid.

`#thread`: number of threads per block.

`size` and `stream` can be ignored in our assignment.

- E.g:


```
AddKernel<<<32, 1024>>(d_c, d_a, d_b);
```


Build-in variable `dim3`

- `dim3` is an integer vector type that can be used in CUDA code.
- Its most common application:
 - pass the grid and block dimensions in a kernel invocation.
- `dim3` has 3 elements: x, y, z
 - C code initialization: `dim3 grid = {512, 512, 1}`
 - C++ code initialization: `dim3 grid(512, 512, 1);`
 - Not all three elements need to be provided.
 - Any element not provided during initialization is initialized to 1, not 0!
- E.g:

```
dim3 block(32); // 32 * 1 * 1
dim3 thread(1024) // 1024 * 1 * 1
AddKernel<<< block, thread>>>(d_c, d_a, d_b);
```

Dim3 example



```
// 1 grid, 4 blocks per grid, 4 threads per block.  
dim3 block(4, 1, 1); //4 blocks per grid  
dim3 thread(4, 1, 1); // 4 threads per block  
addKernel<<<block, thread>>>(d_c, d_a, d_b);
```

Thread index calculation

- 1D grid of 1D blocks



```
//1D * 1D  
threadID = blockDim.x * blockIdx.x + threadIdx.x;
```

- 1D grid of 2D blocks

```
//1D * 2D  
threadID = blockDim.x * blockDim.y * blockIdx.x +  
           blockDim.x * threadIdx.y +  
           threadIdx.x;
```

- 1D grid of 3D blocks

```
//1D * 3D  
threadID = blockDim.x * blockDim.y * blockDim.z * blockIdx.x +  
           (blockDim.x * blockDim.y) * threadIdx.z +  
           blockDim.x * threadIdx.y +  
           threadIdx.x;
```

Thread index calculation

- 2D grid of 1D blocks

```
//2D * 1D
blockID = gridDim.x * blockIdx.y + blockIdx.x;
threadID = blockID * blockDim.x + threadIdx.x
```

- 2D grid of 2D blocks

```
//2D * 2D
blockID = gridDim.x * blockIdx.y + blockIdx.x;
threadID = blockID * (blockDim.x * blockDim.y) +
           blockDim.x * threadIdx.y +
           threadIdx.x;
```

- 2D grid of 3D blocks

```
//2D * 3D
blockID = gridDim.x * blockIdx.y + blockIdx.x;
threadID = blockID * (blockDim.x * blockDim.y * blockDim.z) +
           (blockDim.x * blockDim.y) * threadIdx.z +
           blockDim.x * threadIdx.y +
           threadIdx.x;
```

- 3D grid of 1D blocks

```
//3D * 1D
blockID = gridDim.x * gridDim.y * blockIdx.z +
           gridDim.x * blockIdx.y +
           blockIdx.x;
```

- 3D grid of 2D blocks

```
//3D * 2D
blockID = gridDim.x * gridDim.y * blockIdx.z +
           gridDim.x * blockIdx.y +
           blockIdx.x;
threadID = blockID * (blockDim.x * blockDim.y) +
           blockDim.x * threadIdx.y +
           threadIdx.x;
```

- 3D grid of 3D blocks

```
//3D * 3D
blockID = gridDim.x * gridDim.y * blockIdx.z +
           gridDim.x * blockIdx.y +
           blockIdx.x;
threadID = blockID * (gridDim.x * gridDim.y * gridDim.z) +
           (gridDim.x * gridDim.y) * threadIdx.z +
           gridDim.x * threadIdx.y +
           threadIdx.x;
```

Assignment3:Compiling on CS lab2

```
Compile: nvcc -std=c++11 -arch=compute_52 -code=sm_52  
        main.cu  
        cuda_smith_waterman_skeleton.cu  
        -o cuda_smith_waterman
```

```
Run: ./cuda_smith_waterman <input file> <num of blocks per grid> <number of thread  
per block>
```

Or just use run_cuda.sh bash script.

Assignment3: 2D score -> 1D score

Linear representation of a 2D array is convenient to achieve **Coalesced Memory Access**.

In assignment3, the score matrix is represented as 1D array.

To be fair, the serial code in assignment3 reacts to the changes.

Index transform function.

```
#pragma once
using namespace std;
const int MATCH = 3, MIS = -3, GAP = 2;
int smith_waterman(char *a, char *b, int a_len, int b_len);
// return score of substitution matrix
inline int sub_mat(char x, char y) {
    return x == y ? MATCH : MIS;
}
inline int idx(int x, int y, int n){
    return x * n + y;
}
```

score matrix allocation changes.

```
int smith_waterman(char *a, char *b, int a_len, int b_len) {
    // init score matrix
    int **score = new int*[a_len + 1];
    for (int i = 0; i <= a_len; i++) {
        score[i] = new int[b_len + 1];
        for (int j = 0; j <= b_len; j++) {
            score[i][j] = 0;
        }
    }

    int *score = (int *)malloc(sizeof(int) * (a_len + 1) * (b_len + 1));
    for(int i = 0; i <= a_len; i++){
        for(int j = 0; j <= b_len; j++){
            score[idx(i, j, b_len + 1)] = 0;
        }
    }
}
```

score matrix addressing changes.

```
// main loop
int max_score = 0;
for (int i = 1; i <= a_len; i++) {
    for (int j = 1; j <= b_len; j++) {
        score[idx(i, j, b_len + 1)] = max(0,
            max(score[i - 1][j - 1] + sub_mat(a[i - 1], b[j - 1]),
                max(score[i - 1][j] - GAP,
                    score[i][j - 1] - GAP)));
        max_score = max(max_score, score[i][j]);
    }
}
```

Assignment3: 2D score -> 1D score

A device version of index transform is provided in the `cuda_smith_waterman.h`

```
#pragma once

const int MATCH = 3, MIS = -3, GAP = 2;

int smith_waterman(int blocks_per_grid, int threads_per_block, char *a, char *b, int a_len, int b_len);

inline __device__ int sub_mat(char x, char y) {
    return x == y ? MATCH : MIS;
}

#define GPUErrChk(ans) { utils::GPUAssert((ans), __FILE__, __LINE__); }

namespace utils {

    inline void GPUAssert(cudaError_t code, const char *file, int line, bool abort = true) {
        if (code != cudaSuccess) {
            fprintf(stderr, "GPU assert: %s %s %d\n", cudaGetErrorString(code), file, line);
            if (abort)
                exit(code);
        }
    }

    inline __device__ int dev_idx(int x, int y, int n) {
        return x * n + y;
    }
}
```

Assignment3: Error check helper function

```
#pragma once

const int MATCH = 3, MIS = -3, GAP = 2;

int smith_waterman(int blocks_per_grid, int threads_per_block, char *a, char *b, int a_len, int b_len);

inline __device__ int sub_mat(char x, char y) {
    return x == y ? MATCH : MIS;
}

#define GPUErrChk(ans) { utils::GPUAssert((ans), __FILE__, __LINE__); }

namespace utils {

    inline void GPUAssert(cudaError_t code, const char *file, int line, bool abort = true) {
        if (code != cudaSuccess) {
            fprintf(stderr, "GPU assert: %s %s %d\n", cudaGetErrorString(code), file, line);
            if (abort)
                exit(code);
        }
    }

    inline __device__ int dev_idx(int x, int y, int n) {
        return x * n + y;
    }
}
```

E.g:


```
GPUErrChk(cudaMalloc(&d_A, size));
```





Assignment3: Your task

- Handle memory allocation & deallocation by yourself.
 - Handle memory copy by yourself.
 - Write one or more kernels to do the computation of score matrix.
-
- Note: Using the global memory is enough.

Assignment3: Hints

- **Coalesced memory access.**
 - Improve performance greatly.
 - If the threads in a block are accessing consecutive global memory locations, then all the accesses are combined into a single request(or coalesced) by the hardware.
 - refer to lecture notes:
 - cuda_programming model
 - n-body simulation
- Memory access pattern: Consecutive threads access consecutive memory addresses.






Consecutive threads access consecutive memory addresses.



Transfer anti-diagonals to rows.

1				
5	2			
9	6	3		
13	10	7	4	
14	11	8		
15	12			
16				

align



1	ε	ε	ε
5	2	ε	ε
9	6	3	ε
13	10	7	4
14	11	8	ε
15	12	ε	ε
16	ε	ε	ε



Assignment3: References

- Kernel launch parameters setup:
 - #block is set to the number of SM(streaming multiprocessors) of GPU or the multiples of SMs.
 - CS lab2 machines' GPU has 8 SMs
 - #thread is set to the multiples of 32.
 - E.g: <<<8, 256>>> <<<16, 1024>>> <<<4, 512>>>
- Referential running time:

```
../test/4k.in 8 512  
Max score: 12327  
Elapsed Time: 0.256098949 s  
Driver Time: 0.160557251 s
```

Assignment3: Suggestion

If you have questions about CUDA programming:

- Work hard
- Read your lecture notes
- Read NVIDIA CUDA documents
- Ask Google
- Read previous year's final exam algorithm code on Maximum Flow problem
- Send emails to TAs